

# FORECASTING FOREIGN EXCHANGE RATE

Soukhyada Vaidya  
Roshni Suhandu  
Daksh Gupta  
Shravankumar Hiregoudar

```
library(readr)
library(tseries)
library(forecast)

exchange <- read.csv("E:/MITA/Bf/project/exchange.csv", header = TRUE)
summary(exchange)
```


```
##      i..DATE      BRL.USD      RUB.USD      INR.USD
## 01-01-1998: 1  Min.   :1.122  Min.   : 5.998  Min.   :38.79
## 01-01-1999: 1  1st Qu.:1.789  1st Qu.:27.260  1st Qu.:43.98
## 01-01-2000: 1  Median :2.114  Median :29.203  Median :46.06
## 01-01-2001: 1  Mean    :2.206  Mean    :30.174  Mean    :48.41
## 01-01-2002: 1  3rd Qu.:2.468  3rd Qu.:31.352  3rd Qu.:50.08
## 01-01-2003: 1  Max.    :3.892  Max.    :70.188  Max.    :66.55
## (Other)      :210
##      YUAN.USD
## Min.   :6.052
## 1st Qu.:6.471
## Median :7.805
## Mean    :7.425
## 3rd Qu.:8.279
## Max.    :8.307
##
```

The currencies are: Brazilian Real Russian Ruble Indian rupees Chinese Yuan

```
head (exchange)

##      i..DATE BRL.USD RUB.USD INR.USD YUAN.USD
## 1 01-01-1998 1.1218 6.0100 39.5402 8.2940
## 2 01-02-1998 1.1264 5.9983 38.7942 8.2780
## 3 01-03-1998 1.1322 6.0100 39.4955 8.2771
## 4 01-04-1998 1.1359 6.0449 39.4695 8.2514
## 5 01-05-1998 1.1441 6.1356 40.1854 8.2621
## 6 01-06-1998 1.1532 6.2005 42.0108 8.2816
```

```
plot(exchange$RUB.USD)
```



Create a time series object

```
exchange <- read.csv("E:/MITA/Bf/project/exchange.csv", header = TRUE)

timeseries=ts(exchange[, "RUB.USD"], frequency=12, start=1998)
exchange$RUB.USD=tsclean(timeseries)
```


Decomposition of time series

It calculate the seasonal component using smoothing and then adjust the original series for seasonality. The result is a seasonality adjusted time series

```
count_ts=ts(na.omit(exchange$RUB.USD), frequency=12, start = 1998)

# Making the data into trend,seasonal and reminder
decomp <- stl(count_ts, s.window = "periodic")


# Seasonally adjusting the data
deseasonal_ts <- seasadj(decomp)
plot(decomp)
```



```
# Plotting the data with trend
plot(count_ts, col="gray", main="Exchange rate (Russian Ruble - USD)",
      ylab="currency value", xlab="Time")


# 1 is seasonal, 2 is trend and 3 is reminder
lines(decomp$time.series[,2],col="red")
```

### Exchange rate (Russian Ruble - USD)



```
plot(count_ts, col="grey", main="Exchange rate (Russian Ruble - USD)",  
ylab="currency value", xlab="Time")  
lines(seasadj(decomp),col="red",ylab="Seasonally adjusted")
```

### Exchange rate (Russian Ruble - USD)



## Stationarity test

Augmented Dickey-Fuller test for stationarity/ non-stationarity of data. A large p-value(>0.05) suggests that the time series is non-stationary (Reject H0 hypothesis)

Kwiatkowski-Phillips-Schmidt-Shin test of stationarity. In this case, if p-value are smaller than 5% differencing is required => for KPSS test H0:stationarity; H1:non-stationarity

```
adf<-adf.test(count_ts, alternative = "stationary")

## Warning in adf.test(count_ts, alternative = "stationary"): p-value greater
## than printed p-value

kpss<-kpss.test(count_ts)

## Warning in kpss.test(count_ts): p-value smaller than printed p-value

adf

##
## Augmented Dickey-Fuller Test
##
## data: count_ts
## Dickey-Fuller = 1.3739, Lag order = 5, p-value = 0.99
## alternative hypothesis: stationary

kpss

##
## KPSS Test for Level Stationarity
##
## data: count_ts
## KPSS Level = 1.8503, Truncation lag parameter = 4, p-value = 0.01
```


Autocorrelation is the linear correlation of a signal with itself at two different points in time, ACF (autocorrelation function) is just such correlation as a function of the lag h between two points of time. It correlates with itself through time.

PACF (partial autocorrelation function) is essentially the autocorrelation of a signal with itself at different points in time, with linear dependency with that signal at shorter lags removed, as a function of lag between points of time.

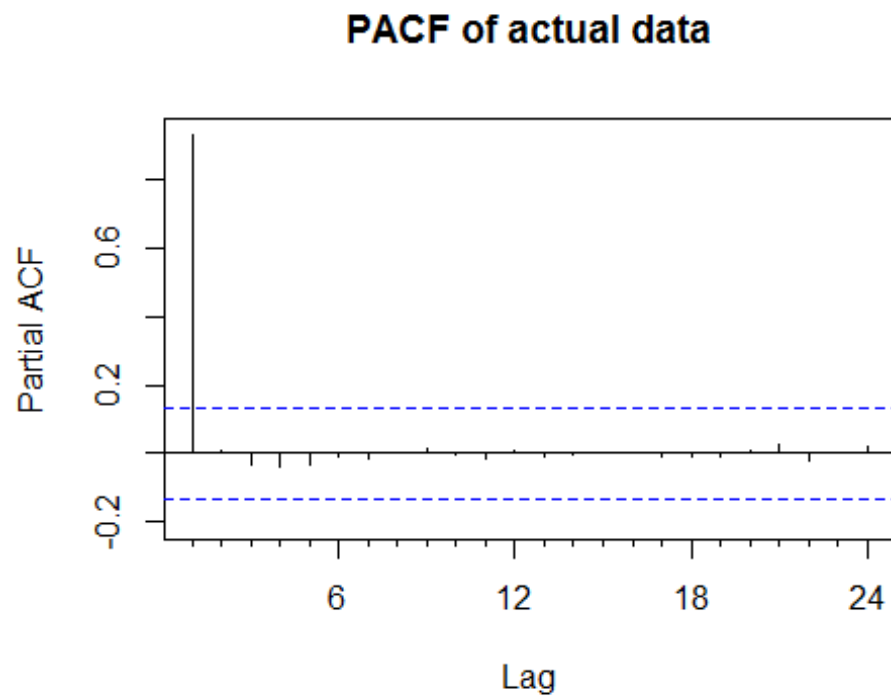
augmented Dickey-Fuller test (ADF) Null Hypothesis (H0): If failed to be rejected, it suggests the time series has a unit root, meaning it is non-stationary. It has some time dependent structure. Alternate Hypothesis (H1): The null hypothesis is rejected; it suggests the time series does not have a unit root, meaning it is stationary. p-value > 0.05: Fail to reject the null hypothesis (H0), the data has a unit root and is non-stationary. p-value <= 0.05: Reject the null hypothesis (H0), the data does not have a unit root and is stationary.

Kwiatkowski-Phillips-Schmidt-Shin (KPSS) test The null hypothesis (H0) for the test is that the data is stationary. The alternate hypothesis (H1) for the test is that the data is not stationary. p-value > 0.05: Reject H0 p-value <= 0.05: Accept H0

```
Acf(count_ts, main='ACF of actual data')
```



```
Pacf(count_ts, main='PACF of actual data')
```



Results of Adf and Kpss test showed that the time series is not stationary.

```
ndiffs(count_ts)
```

```
## [1] 2
```

So, The data gets stationary after 2 differences.

Differencing the data  $d = 1$ ;

```
# deseasonal_ts <- seasadj(decomp) is seasonally adjusted data  
ts_d1=diff(deseasonal_ts, differences=1)
```

```
#stationary test for d=1
```

```
adfd1<-adf.test(ts_d1, alternative = "stationary")
```

```
adfd1
```

```
##
```

```
## Augmented Dickey-Fuller Test
```

```
##
```

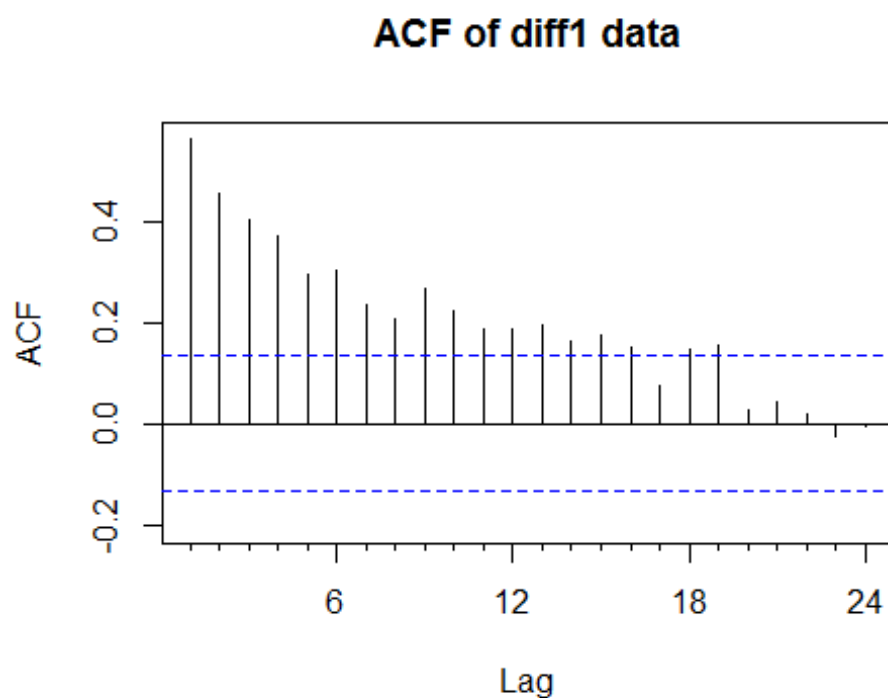
```
## data: ts_d1
```

```
## Dickey-Fuller = -1.9613, Lag order = 5, p-value = 0.5922
```

```
## alternative hypothesis: stationary
```


Even now  $p > 0.05$  suggests that the time series is non-stationary (Reject  $H_0$  hypothesis)

```
Acf(ts_d1, main='ACF of diff1 data')
```



```
Pacf(ts_d1, main='PACF of diff1 data')
```

### PACF of diff1 data



Differencing the data  $d = 2$ ;

```
# deseasonal_ts <- seasadj(decomp) is seasonally adjusted data
ts_d2=diff(ts_d1, differences=1)

#stationary test for d=2
adfd2<-adf.test(ts_d2, alternative = "stationary")

## Warning in adf.test(ts_d2, alternative = "stationary"): p-value smaller
## than printed p-value

adfd2


##
## Augmented Dickey-Fuller Test
##
## data: ts_d2
## Dickey-Fuller = -7.932, Lag order = 5, p-value = 0.01
## alternative hypothesis: stationary
```

Now,  $p < 0.05$  suggests that the time series is stationary (Accept  $H_0$  hypothesis)

```
Acf(ts_d2, main='ACF of diff2 data')
```




**ACF of diff2 data**



```
Pacf(ts_d2, main='PACF of diff1 data')
```

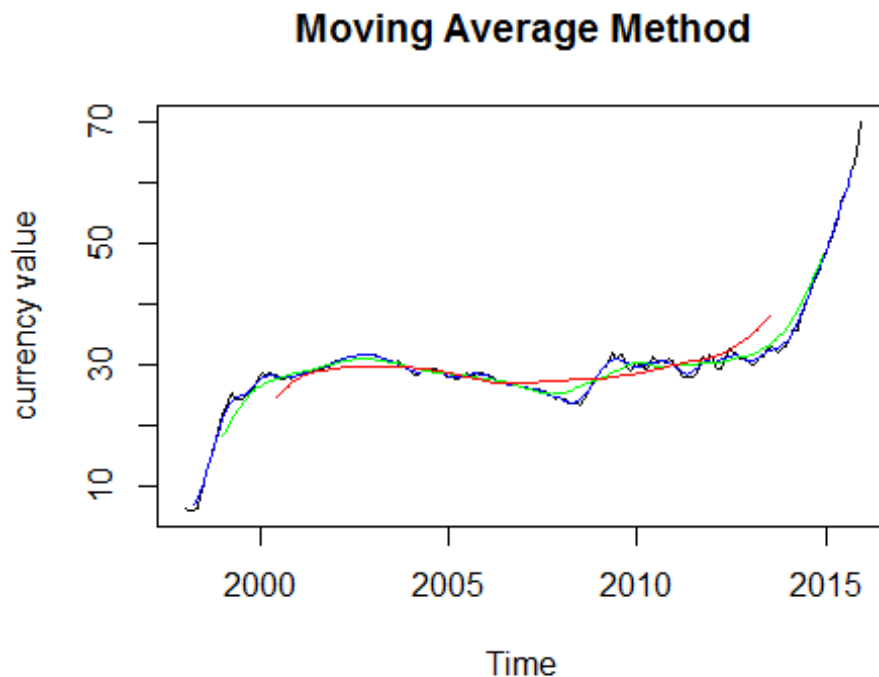
**PACF of diff1 data**



After a time series has been stationarized by differencing, the next step in fitting an ARIMA model is to determine whether AR or MA terms are needed to correct any autocorrelation that remains in the differenced series. By looking at the autocorrelation function (ACF) and partial autocorrelation (PACF) plots of the differenced series, you can tentatively identify the numbers of AR and/or MA terms that are needed. ACF plot: it is merely a bar chart of the coefficients of correlation between a time series and lags of itself. The PACF plot is a plot of the partial correlation coefficients between the series and lags of itself.

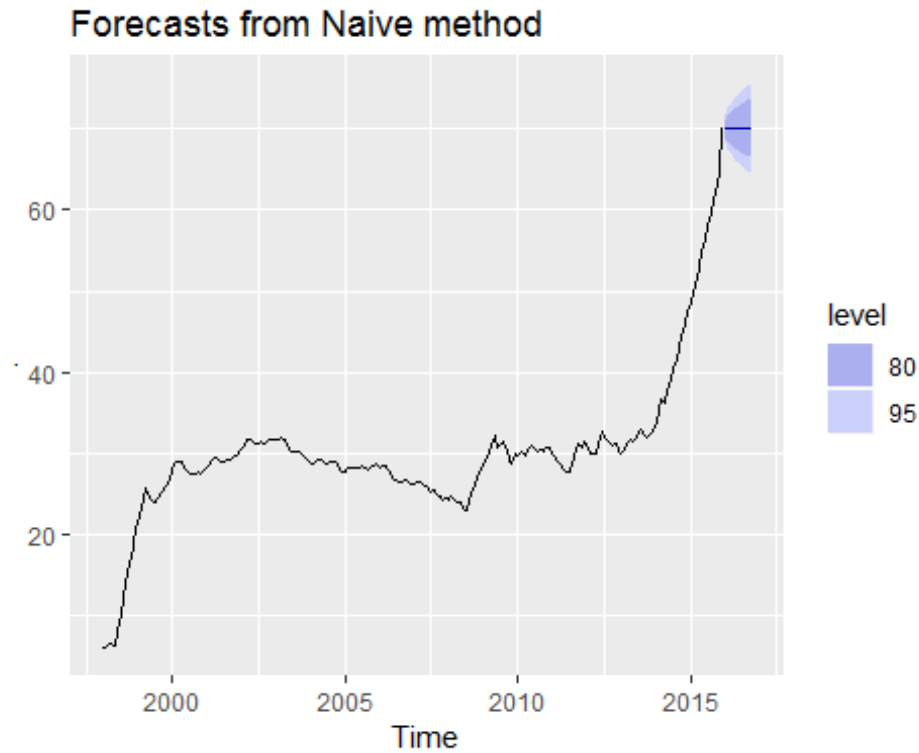
Moving average of 7 for the stationary data

```
plot(count_ts, main="Moving Average Method", ylab="currency value",
      xlab="Time")
lines(ma(count_ts,7),col="blue")
lines(ma(count_ts,25), col="green")
lines(ma(count_ts,59), col="red")
```



Naive forecasts of seasonally adjusted data


```
fit <- stl(count_ts, t.window=30, s.window="periodic", robust=TRUE)
fit %>% seasadj() %>% naive() %>%
# plotting the graph without seasonality
autoplot()
```



```
library(ggplot2)

#divide the data into training and testing
train <- deseasonal_ts[1:150]
test <- deseasonal_ts[150:216]

fit1 <- meanf(count_ts,h=67)
fit2 <- rwf(count_ts,h=67)
fit3 <- snaive(count_ts,h=67)
autoplot(window(count_ts, start=1998)) +
  autolayer(fit1, series="Mean", PI=FALSE) +
  autolayer(fit2, series="Naïve", PI=FALSE) +
  autolayer(fit3, series="Seasonal naïve", PI=FALSE) +
  ylab("currency rates") + xlab("year") +
  ggtitle("Forecasts for currency exchange rate") +
  guides(colour=guide_legend(title="Forecast"))
```



```
accuracy(fit1, test)
```

```
##               ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 1.842341e-16  8.896925  4.914027 -13.50684  23.74663   8.449778
## Test set     7.221067e+00 12.891380  7.566082  14.23475  15.44650  13.010045
##               ACF1
## Training set 0.9314482
## Test set     NA
```


```
accuracy(fit2, test)
```

```
##               ME      RMSE      MAE      MPE      MAPE
## Training set  0.2972487  0.8824197  0.5815569   1.044885   2.068504
## Test set     -32.9360934 34.6241264 32.9360934 -100.450309 100.450309
##               MASE      ACF1
## Training set  1.00000 0.549123
## Test set     56.63434   NA
```


```
accuracy(fit3, test)
```

```
##               ME      RMSE      MAE      MPE      MAPE      MASE
## Training set  2.747478  6.327891  3.846758   7.482198  11.52544   6.614585
## Test set     -19.879916 23.581694 21.746094 -63.377021  66.38235  37.392890
##               ACF1
## Training set 0.9452067
## Test set     NA
```

```
## Fit a model by using Holt method and find the fitted values (estimation of
original values)
fc <- holt(count_ts, h=67)
autoplot(count_ts) +
  autolayer(fc, series="Holt's method", PI=FALSE) +
  ggtitle("Forecasts from Holt's method") + xlab("Time") +
  ylab("Currency Value") + guides(colour=guide_legend(title="Forecast"))
```



```
plot(fitted(fc))
lines(fitted(fc), col = "blue", type = "o")
```



**fitted(fc)**

##	Jan	Feb	Mar	Apr	May	Jun	Jul
## 1998	3.536353	5.872186	5.642468	5.803458	5.936461	6.108069	9.317140
## 1999	23.104476	23.412559	24.638446	25.239072	26.694645	25.101015	24.614057
## 2000	27.409035	29.049309	29.449095	28.743394	28.790663	28.309539	28.193766
## 2001	28.120343	28.592884	28.840408	28.853197	29.037568	29.203446	29.276477
## 2002	30.416573	30.889087	31.153413	31.374475	31.403986	31.422300	31.594141
## 2003	31.919892	31.866792	31.606874	31.312828	31.025659	30.684093	30.151131
## 2004	29.127857	28.404019	28.123282	28.311000	28.676828	29.066595	29.090509
## 2005	27.500684	27.719589	27.835857	27.423635	27.759576	27.972438	28.770183
## 2006	28.893507	28.236420	28.072963	27.647179	27.281334	26.689852	26.773280
## 2007	26.089994	26.483118	26.223012	25.964247	25.627962	25.714869	25.881057
## 2008	24.410545	24.353889	24.437525	23.363056	23.231185	23.629935	23.557423
## 2009	29.012954	29.597989	30.045806	30.444570	31.157613	32.869919	31.289627
## 2010	30.139576	29.801105	30.281237	29.397204	28.935845	30.835127	31.787468
## 2011	30.862074	29.891147	28.767051	27.800121	27.547433	27.550702	27.792371
## 2012	32.001767	31.529423	29.306106	28.897218	29.245568	31.400276	33.924247
## 2013	30.443889	29.877326	29.958127	30.912345	31.596443	31.509328	32.845778
## 2014	33.041164	34.294201	36.221031	36.992035	36.019515	38.579598	40.726595
## 2015	48.816353	50.112463	50.945946	52.256264	53.620012	56.268083	58.581657
##	Aug	Sep	Oct	Nov	Dec		
## 1998	11.657845	14.014312	16.245620	18.615669	18.665528		
## 1999	24.418226	24.943943	25.911023	26.124907	26.765129		
## 2000	27.657337	27.576442	27.739531	27.865089	27.814998		
## 2001	29.362780	29.479581	29.552139	29.672078	30.007151		
## 2002	31.674031	31.679368	31.738927	31.789256	31.927183		

```
## 2003 30.121604 30.225824 30.616172 29.931444 29.558147
## 2004 29.139399 29.307376 29.265293 29.014037 28.344180
## 2005 28.926765 28.524878 28.355072 28.639316 28.905860
## 2006 26.743803 26.586082 26.658625 26.856633 26.480291
## 2007 25.367777 25.572241 25.145342 24.607440 24.127359
## 2008 23.183237 24.464455 25.912213 27.265719 28.171810
## 2009 31.815090 31.927662 30.515530 28.775423 28.360503
## 2010 30.721933 30.341692 30.940718 30.209883 31.198748
## 2011 27.788064 29.057827 31.711756 32.066781 31.132782
## 2012 32.983167 32.022761 31.253052 30.865959 31.348803
## 2013 33.243699 33.433079 32.662691 31.904958 32.881291
## 2014 41.908742 43.035599 44.714199 45.812306 47.527169
## 2015 60.010293 61.412163 63.385228 64.801326 66.744771
```

```
fc$model
```

```
## Holt's method
##
## Call:
## holt(y = count_ts, h = 67)
##
## Smoothing parameters:
##   alpha = 0.9999
##   beta  = 0.4063
##
## Initial states:
##   l = 5.0579
##   b = -1.5215
##
## sigma: 0.6822
##
##      AIC      AICc      BIC
## 1001.783 1002.069 1018.659
```


```
accuracy(fc, test)
```

```
##              ME      RMSE      MAE      MPE      MAPE
## Training set  0.05298339 0.6758088 0.4394268 0.4846092 1.834466
## Test set     -139.28644808 148.6694513 139.2864481 -371.6019886 371.601989
##              MASE      ACF1
## Training set  0.755604 0.08076586
## Test set     239.506133      NA
```

```
## Fit a model by using Holt-Winters' method with additive seasonality and
## find the fitted values (estimation of original values)
```

```
fit1 <- hw(count_ts, seasonal="additive")
fit2 <- hw(count_ts, seasonal="multiplicative")
autoplot(count_ts) +
  autolayer(fit1, series="HW additive forecasts", PI=FALSE) +
  autolayer(fit2, series="HW multiplicative forecasts", PI=FALSE) +
```

```
xlab("Year") +
ylab("Currency values") +
ggtitle("Holt-Winters' method") +
guides(colour=guide_legend(title="Forecast"))
```



```
fit1$model

## Holt-Winters' additive method
##
## Call:
## hw(y = count_ts, seasonal = "additive")
##
## Smoothing parameters:
##   alpha = 0.9367
##   beta  = 0.1324
##   gamma = 0.0627
##
## Initial states:
##   l = 4.1638
##   b = 0.9608
##   s = -0.1025 -0.1144 -0.2646 -0.306 -0.2623 -0.019
##        0.0483 0.3358 0.0594 -0.0092 0.0261 0.6084
##
## sigma: 0.7346
##
##      AIC      AICc      BIC
## 1045.170 1048.261 1102.549
```



```

fit2$model

## Holt-Winters' multiplicative method
##
## Call:
## hw(y = count_ts, seasonal = "multiplicative")
##
## Smoothing parameters:
##   alpha = 0.4989
##   beta  = 0.2011
##   gamma = 1e-04
##
## Initial states:
##   l = 4.4614
##   b = 0.3362
##   s = 1.011 0.999 0.9999 1.0075 1.0023 1.0021
##       1.0032 0.995 0.9877 0.9908 0.9975 1.0039
##
## sigma: 0.049
##
##      AIC      AICc      BIC
## 1323.179 1326.270 1380.559

# since AIC of additive is less. It is good

```

Fit the ARIMA model.

ARIMA stands for auto-regressive integrated moving average and is specified by these three order parameters: (p, d, q).

AR(p), is referring to the use of past values in the regression equation for the series Y. The auto-regressive parameter p specifies the number of lags used in the model.

I(d) The d represents the degree of differencing

MA(q) component represents the error of the model as a combination of previous error terms. The order q determines the number of terms to include in the model

auto.arima() command choose the best ARIMA model

```

#fitting the model
model1<-auto.arima(deseasonal_ts, seasonal = F)
model1

## Series: deseasonal_ts
## ARIMA(1,2,1)
##
## Coefficients:
##      ar1      ma1
##    0.3208 -0.8394
## s.e. 0.0975 0.0587
##

```

```
## sigma^2 estimated as 0.3857: log likelihood=-201.07
## AIC=408.14 AICc=408.26 BIC=418.24
```

Evaluate the model for forecast.

we would expect no significant autocorrelations present in the model residuals.

```
forecast(model1,h = 20)
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Jan 2016	72.94141	72.14546	73.73736	71.72410	74.15871
## Feb 2016	75.55951	74.13684	76.98219	73.38372	77.73530
## Mar 2016	78.00299	75.98585	80.02014	74.91804	81.08795
## Apr 2016	80.39045	77.79180	82.98910	76.41616	84.36474
## May 2016	82.75994	79.57984	85.94003	77.89641	87.62346
## Jun 2016	85.12365	81.35444	88.89287	79.35914	90.88817
## Jul 2016	87.48552	83.11503	91.85601	80.80143	94.16961
## Aug 2016	89.84679	84.86032	94.83327	82.22064	97.47295
## Sep 2016	92.20788	86.58930	97.82646	83.61500	100.80075
## Oct 2016	94.56890	88.30134	100.83646	84.98349	104.15431
## Nov 2016	96.92990	89.99610	103.86370	86.32557	107.53424
## Dec 2016	99.29090	91.67349	106.90831	87.64108	110.94072
## Jan 2017	101.65189	93.33352	109.97027	88.93004	114.37375
## Feb 2017	104.01289	94.97632	113.04946	90.19265	117.83313
## Mar 2017	106.37388	96.60205	116.14571	91.42916	121.31861
## Apr 2017	108.73488	98.21093	119.25882	92.63989	124.82986
## May 2017	111.09587	99.80318	122.38856	93.82520	128.36654
## Jun 2017	113.45686	101.37905	125.53468	94.98544	131.92829
## Jul 2017	115.81786	102.93878	128.69694	96.12100	135.51471
## Aug 2017	118.17885	104.48261	131.87510	97.23225	139.12545

Lets divide the data and find the accuracy

Take the training set as all the values except the last 10 (test data) and find the accuracy.

```
#divide the data into training and testing
```

```
train <- deseasonal_ts[1:150]
test <- deseasonal_ts[150:216]
```

```
# find the best fit
```

```
fit <- auto.arima(train)
fit
```

```
## Series: train
```

```
## ARIMA(1,2,2)
```

```
##
```

```
## Coefficients:
```

```
##          ar1          ma1          ma2
```

```
##          0.8254      -1.4065      0.4175
```


```
## s.e.      0.1335      0.1599      0.1332
```

```
##
```

```
## sigma^2 estimated as 0.2879: log likelihood=-117.1
## AIC=242.2   AICc=242.48   BIC=254.19

# fit the model and predict the next 67 values
testtest<- forecast(fit,h = 67)
plot(testtest)
```

### Forecasts from ARIMA(1,2,2)



```
# find the accuracy of fitted and actual
accuracy(testtest,test)
```

	ME	RMSE	MAE	MPE	MAPE
## Training set	-0.03513822	0.5275399	0.3566533	0.06092997	1.535543
## Test set	-0.38850643	8.5036860	6.7091207	-5.95947078	16.961706
	MASE	ACF1			
## Training set	0.796456	-0.03842399			
## Test set	14.982390	NA			